



Développeur d'Intelligence Artificielle Appliquée

Cours #6

www.impactia.org

Structure de la formation


- **#1: Introduction**
- **#2: Vision**
- **#3: Vision**
- **#4: Vision**
- **#5: Renforcement**
- **#6: Renforcement**
- #7: Renforcement
- #8: Langage
- #9: Langage
- #10: Langage
- #11: Outillage
- #12: Génération d'images
- #13: Génération d'images
- #14: Projet
- #15: Projet


Semaine dernière : Cours #4


- Théorie
 - Introduction à l'apprentissage par renforcement
 - Que peut-on faire avec ?
 - Pourquoi est-ce difficile ?
- Pratique
 - L'environnement droneRL

Aujourd'hui : Cours #5

- Théorie
 - Q-learning
 - Deep Q-Network (DQN)
- Pratique
 - Entrainement d'agents droneRL
 - Premières participations au Challenge droneRL







But :

Maximiser la somme des *récompenses* ("rewards")
jusqu'à la fin de l'épisode

The Promise of RL


Specify **what** to do, but not **how** to do it

- Through the reward function
- Learning “fills in the details”


Better final solutions


- Based of actual experiences, not programmer assumptions

Less (human) time needed for a good solution




L'environnement *DeliveryDrones*





```
self.env_params = {  
    'drone_density': 0.05,  
    'n_drones': 3,  
    'pickup_reward': 0,  
    'delivery_reward': 1,  
    'crash_reward': -1,  
    'charge_reward': -0.1,  
    'discharge': 10,  
    'charge': 20,  
    'packets_factor': 3,  
    'dropzones_factor': 2,  
    'stations_factor': 2,  
    'skyscrapers_factor': 3,  
    'rgb_render_rescale': 1.0  
}
```



Mise en pratique

L'environnement DeliveryDrones

[https://colab.research.google.com/drive/
18R9mQPDRv-rhVons1IEPjmCCiP7sjhsz](https://colab.research.google.com/drive/18R9mQPDRv-rhVons1IEPjmCCiP7sjhsz)



Comment “apprendre un comportement” ? *Q-learning*

Apprendre un comportement




Une information à disposition



5 actions possibles - que faire ?

LEFT DOWN RIGHT UP STAY



But :

Maximiser la somme des *récompenses* ("rewards")
jusqu'à la fin de l'épisode

Actions

Observation		LEFT	DOWN	RIGHT	UP	STAY
	→	1.7	1.4	2.6	0.86	1.5
	↘	1.3	1.6	2.4	0.78	1.5
	↑	1.9	1.6	1.4	2.5	1.7
	↗	1.4	0.58	1.8	1.6	1.6
	↖	2.2	1.3	0.97	1.6	1.5
	↙	2.5	2	1.3	0.9	1.5
	↓	1.8	2.9	1.7	1.2	2.1
	←	2.6	1.5	2.1	1.2	2.2

La “solution” d'un environnement

Si on connaît les valeurs de ce tableau,
il est facile de suivre un comportement optimal !

		Actions				
		LEFT	DOWN	RIGHT	UP	STAY
Observation	→	1.7	1.4	2.6	0.86	1.5
	↘	1.3	1.6	2.4	0.78	1.5
	↑	1.9	1.6	1.4	2.5	1.7
	↗	1.4	0.58	1.8	1.6	1.6
	↖	2.2	1.3	0.97	1.6	1.5
	↙	2.5	2	1.3	0.9	1.5
	↓	1.8	2.9	1.7	1.2	2.1
	←	2.6	1.5	2.1	1.2	2.2

“Q-table”

La “solution” d'un environnement


$Q(\text{observation}, \text{action}) = \text{somme des récompenses attendues}$

		Actions				
		LEFT	DOWN	RIGHT	UP	STAY
Observation	→	1.7	1.4	2.6	0.86	1.5
	↘	1.3	1.6	2.4	0.78	1.5
	↑	1.9	1.6	1.4	2.5	1.7
	↗	1.4	0.58	1.8	1.6	1.6
	↖	2.2	1.3	0.97	1.6	1.5
	↙	2.5	2	1.3	0.9	1.5
	↓	1.8	2.9	1.7	1.2	2.1
	←	2.6	1.5	2.1	1.2	2.2

“Q-table”


Comment apprendre la Q-table ?

Dynamic Programming




Comment apprendre la Q-table ?

Monte-Carlo



Comment apprendre la Q-table ?

Temporal-Difference



Q-learning

Actions

	LEFT	DOWN	RIGHT	UP	STAY
→	1.7	1.4	2.6	0.86	1.5
↓	1.3	1.6	2.4	0.78	1.5
↑	1.9	1.6	1.4	2.5	1.7
↗	1.4	0.58	1.8	1.6	1.6
↖	2.2	1.3	0.97	1.6	1.5
↙	2.5	2	1.3	0.9	1.5
↘	1.8	2.9	1.7	1.2	2.1
↔	2.6	1.5	2.1	1.2	2.2



“Q-table”

Somme attendue des récompenses



Actions

	LEFT	DOWN	RIGHT	UP	STAY
→	1.7	1.4	2.6	0.86	1.5
↓	1.3	1.6	2.4	0.78	1.5
↑	1.9	1.6	1.4	2.5	1.7
↗	1.4	0.58	1.8	1.6	1.6
↖	2.2	1.3	0.97	1.6	1.5
↙	2.5	2	1.3	0.9	1.5
↘	1.8	2.9	1.7	1.2	2.1
↔	2.6	1.5	2.1	1.2	2.2



“Q-table”

Somme attendue des récompenses



Observation

Actions


	LEFT	DOWN	RIGHT	UP	STAY
→	1.7	1.4	2.6	0.86	1.5
↓	1.3	1.6	2.4	0.78	1.5
↑	1.9	1.6	1.4	2.5	1.7
↗	1.4	0.58	1.8	1.6	1.6
↖	2.2	1.3	0.97	1.6	1.5
↙	2.5	2	1.3	0.9	1.5
↘	1.8	2.9	1.7	1.2	2.1
←	2.6	1.5	2.1	1.2	2.2

“Q-table”

Somme attendue des récompenses

On sélectionne l'action qui maximise la somme attendue des récompenses

$$a = \operatorname{argmax}_a Q(s, a)$$



Observation

Actions

	LEFT	DOWN	RIGHT	UP	STAY
→	1.7	1.4	2.6	0.86	1.5
↓	1.3	1.6	2.4	0.78	1.5
↑	1.9	1.6	1.4	2.5	1.7
↗	1.4	0.58	1.8	1.6	1.6
↖	2.2	1.3	0.97	1.6	1.5
↙	2.5	2	1.3	0.9	1.5
↘	1.8	2.9	1.7	1.2	2.1
←	2.6	1.5	2.1	1.2	2.2

“Q-table”

Somme attendue des récompenses

On sélectionne l'action qui maximise la somme attendue des récompenses

$$a = \operatorname{argmax}_a Q(s, a)$$



Observation


		Actions				
		LEFT	DOWN	RIGHT	UP	STAY
Observation	→	1.7	1.4	2.6	0.86	1.5
	↓	1.3	1.6	2.4	0.78	1.5
	↑	1.9	1.6	1.4	2.5	1.7
	↗	1.4	0.58	1.8	1.6	1.6
	↖	2.2	1.3	0.97	1.6	1.5
	↙	2.5	2	1.3	0.9	1.5
	↓	1.8	2.9	1.7	1.2	2.1
	←	2.6	1.5	2.1	1.2	2.2


“Q-table”


Somme attendue des récompenses

On sélectionne l'action qui maximise la somme attendue des récompenses


$$a = \operatorname{argmax}_a Q(s, a)$$








Action at time t
eg “LEFT”



$$= \sum r_{t+1} + r_{t+2} + \dots + r_T$$



Q-learning

“Equation de Bellman”

$$Q_{n+1}(s_t, a_t) = r_{t+1} + \gamma \max_a Q_n(s_{t+1}, a)$$

	LEFT	DOWN	RIGHT	UP	STAY
→	1.7	1.4	2.6	0.86	1.5
↘	1.3	1.6	2.4	0.78	1.5
↑	1.9	1.6	1.4	2.5	1.7
↗	1.4	0.58	1.8	1.6	1.6
↖	2.2	1.3	0.97	1.6	1.5
↙	2.5	2	1.3	0.9	1.5
↓	1.8	2.9	1.7	1.2	2.1
←	2.6	1.5	2.1	1.2	2.2



Q-learning

		Actions				
		LEFT	DOWN	RIGHT	UP	STAY
States	→	1.7	1.4	2.6	0.86	1.5
	↘	1.3	1.6	2.4	0.78	1.5
	↑	1.9	1.6	1.4	2.5	1.7
	↗	1.4	0.58	1.8	1.6	1.6
	↖	2.2	1.3	0.97	1.6	1.5
	↙	2.5	2	1.3	0.9	1.5
	↓	1.8	2.9	1.7	1.2	2.1
	↔	2.6	1.5	2.1	1.2	2.2


$$a = \operatorname{argmax}_a Q(s_t, a)$$



Equation de Bellman




“Explorer” ou “exploiter” ?



Explorer... parfois !

Méthode “epsilon-greedy”



Q-learning

States

	Actions				
	LEFT	DOWN	RIGHT	UP	STAY
→	1.7	1.4	2.6	0.86	1.5
↘	1.3	1.6	2.4	0.78	1.5
↑	1.9	1.6	1.4	2.5	1.7
↗	1.4	0.58	1.8	1.6	1.6
↖	2.2	1.3	0.97	1.6	1.5
↙	2.5	2	1.3	0.9	1.5
↓	1.8	2.9	1.7	1.2	2.1
←	2.6	1.5	2.1	1.2	2.2

$$a = \operatorname{argmax}_a Q(s_t, a)$$

+
epsilon-greedy



Bellman Equation



Q-learning



Fonctionne dans tous les cas !*

States

	Actions				
	LEFT	DOWN	RIGHT	UP	STAY
→	1.7	1.4	2.6	0.86	1.5
↘	1.3	1.6	2.4	0.78	1.5
↑	1.9	1.6	1.4	2.5	1.7
↗	1.4	0.58	1.8	1.6	1.6
↖	2.2	1.3	0.97	1.6	1.5
↙	2.5	2	1.3	0.9	1.5
↓	1.8	2.9	1.7	1.2	2.1
←	2.6	1.5	2.1	1.2	2.2

$$a = \operatorname{argmax}_a Q(s_t, a)$$

+
epsilon-greedy



Bellman Equation



Q-learning



Fonctionne dans tous les cas !*

States

	Actions				
	LEFT	DOWN	RIGHT	UP	STAY
→	1.7	1.4	2.6	0.86	1.5
↘	1.3	1.6	2.4	0.78	1.5
↑	1.9	1.6	1.4	2.5	1.7
↗	1.4	0.58	1.8	1.6	1.6
↖	2.2	1.3	0.97	1.6	1.5
↙	2.5	2	1.3	0.9	1.5
↓	1.8	2.9	1.7	1.2	2.1
←	2.6	1.5	2.1	1.2	2.2

$$a = \operatorname{argmax}_a Q(s_t, a)$$

+
epsilon-greedy



Bellman Equation




* peut prendre un temps infini 😅

Bellman Equation

$$\underbrace{Q_{n+1}(s_t, a_t)}_{\text{new Q-value}} = r_{t+1} + \gamma \max_a Q_n(s_{t+1}, a)$$

	LEFT	DOWN	RIGHT	UP	STAY
→	1.7	1.4	2.6	0.86	1.5
↓	1.3	1.6	2.4	0.78	1.5
↑	1.9	1.6	1.4	2.5	1.7
↗	1.4	0.58	1.8	1.6	1.6
↖	2.2	1.3	0.97	1.6	1.5
↙	2.5	2	1.3	0.9	1.5
↓	1.8	2.9	1.7	1.2	2.1
↖	2.6	1.5	2.1	1.2	2.2



Bellman Equation


$$Q_{n+1}(s_t, a_t) = r_{t+1} + \underbrace{\gamma \max_a Q_n(s_{t+1}, a)}_{\text{new Q-value}}$$

	LEFT	DOWN	RIGHT	UP	STAY
→	1.7	1.4	2.6	0.86	1.5
↓	1.3	1.6	2.4	0.78	1.5
↑	1.9	1.6	1.4	2.5	1.7
↗	1.4	0.58	1.8	1.6	1.6
↖	2.2	1.3	0.97	1.6	1.5
↙	2.5	2	1.3	0.9	1.5
↓	1.8	2.9	1.7	1.2	2.1
↖	2.6	1.5	2.1	1.2	2.2




Bellman Equation

$$Q_{n+1}(s_t, a_t) = r_{t+1} + \gamma \max_a Q_n(s_{t+1}, a)$$



Discount factor
("facteur d'actualisation")

Learning rate



$$\underbrace{Q_{n+1}(s_t, a_t)}_{\text{new Q-value}} = (1 - \alpha)Q_n(s_t, a_t) + \alpha(r_{t+1} + \gamma \max_a Q_n(s_{t+1}, a))$$

```
alpha = 0    # No update  
alpha = 1    # Full update  
alpha = 0.5 # Mean between old/new
```


Mise en pratique

Introduction à Q-learning

[https://colab.research.google.com/drive/
1KSei3ZdjNyyUYsPMqTyl4G2rm9v-P2E](https://colab.research.google.com/drive/1KSei3ZdjNyyUYsPMqTyl4G2rm9v-P2E)



Q-table size




```
{2: 'target: ↓, lidar: →, ↗, ↑, ↘'}
```

```
Dict(target_dir:Discrete(8), lidar:MultiBinary(8))
```

```
{'target_dir': 2, 'lidar': [0, 0, 0, 0, 1, 1, 1, 1]}
```

$$n_{\text{states}} = 8 * (2^{**8}) = 2048$$

Q-table size



{2: 'target: ↓, lidar: →, ↗, ↑, ↘'}



Dict(target_dir:Discrete(8), lidar:MultiBinary(8))




{'target_dir': 2, 'lidar': [0, 0, 0, 0, 1, 1, 1, 1]}



n_states = 8 * (2**8) = 2048


Deep Q-network (DQN)




“Q-network”

$$\delta = Q_n(s_t, a_t) - (r_{t+1} + \gamma \max Q_n(s_{t+1}, a))$$


δ is the loss to minimize




Deep Q-network (DQN)



Dense layer



Convolutional layer



Timestep: 1, action: left


Reward: 1.0

Timestep: 2, action: right

Reward: 1.0

Timestep: 3, action: left

Reward: 1.0



Timestep: 4, action: left

Reward: 1.0

Timestep: 5, action: left

Reward: 1.0

Timestep: 6, action: right

Reward: 0.0



