



# Développeur d'Intelligence Artificielle Appliquée

*Cours #7*

[www.impactia.org](http://www.impactia.org)

# Structure de la formation


- **#1: Introduction**
- **#2: Vision**
- **#3: Vision**
- **#4: Vision**
- **#5: Renforcement**
- **#6: Renforcement**
- **#7: Renforcement**
- #8: Langage
- #9: Langage
- #10: Langage
- #11: Génération d'images
- #12: Génération d'images
- #13: Génération d'images
- #14: Projet
- #15: Projet


# Semaine dernière : Cours #6


- Théorie
  - L'environnement droneRL
  - Q-learning
- Pratique
  - Entrainement des premiers agents droneRL

# Aujourd'hui : Cours #5

- Théorie
  - Deep Q-Network (DQN)
  - Apprentissage priorisé
- Pratique
  - Entrainement d'agents droneRL
  - Challenge droneRL ! 










**But :**

Maximiser la somme des *récompenses* ("rewards")  
jusqu'à la fin de l'épisode



# L'environnement *DeliveryDrones*





```

# paramètres de l'environnement
self.env_params = {
    'pickup_reward': 0,
    'delivery_reward': 1,
    'crash_reward': -1,
    'charge_reward': -0.1,
    'discharge': 10,
    'charge': 20,
    'drone_density': 0.05,
    'packets_factor': 3,
    'dropzones_factor': 2,
    'stations_factor': 2,
    'skyscrapers_factor': 3,
    'rgb_render_rescale': 1.0
}
# wrapper
env = WindowedGridView(env, radius=3)

```

Paramètres utilisés pour l'évaluation sur Alcrowd !



# Comment “apprendre un comportement” ?

# Apprendre un comportement



Une information à disposition



5 actions possibles - que faire ?

LEFT

DOWN

RIGHT

UP

STAY


# La “solution” d'un environnement

Si on connaît les valeurs de ce tableau,  
il est facile de suivre un comportement optimal !

		Actions				
		LEFT	DOWN	RIGHT	UP	STAY
Observation	→	1.7	1.4	2.6	0.86	1.5
	↘	1.3	1.6	2.4	0.78	1.5
	↑	1.9	1.6	1.4	2.5	1.7
	↗	1.4	0.58	1.8	1.6	1.6
	↖	2.2	1.3	0.97	1.6	1.5
	↙	2.5	2	1.3	0.9	1.5
	↓	1.8	2.9	1.7	1.2	2.1
	←	2.6	1.5	2.1	1.2	2.2

“Q-table”

# Comment apprendre la Q-table ?



# Q-learning

## Actions

	LEFT	DOWN	RIGHT	UP	STAY
→	1.7	1.4	2.6	0.86	1.5
↓	1.3	1.6	2.4	0.78	1.5
↑	1.9	1.6	1.4	2.5	1.7
↗	1.4	0.58	1.8	1.6	1.6
↖	2.2	1.3	0.97	1.6	1.5
↙	2.5	2	1.3	0.9	1.5
↘	1.8	2.9	1.7	1.2	2.1
↔	2.6	1.5	2.1	1.2	2.2



“Q-table”

Somme attendue des récompenses



## Actions

	LEFT	DOWN	RIGHT	UP	STAY
→	1.7	1.4	2.6	0.86	1.5
↓	1.3	1.6	2.4	0.78	1.5
↑	1.9	1.6	1.4	2.5	1.7
↗	1.4	0.58	1.8	1.6	1.6
↖	2.2	1.3	0.97	1.6	1.5
↙	2.5	2	1.3	0.9	1.5
↘	1.8	2.9	1.7	1.2	2.1
↔	2.6	1.5	2.1	1.2	2.2



“Q-table”

Somme attendue des récompenses



**Observation**

**Actions**


	LEFT	DOWN	RIGHT	UP	STAY
→	1.7	1.4	2.6	0.86	1.5
↓	1.3	1.6	2.4	0.78	1.5
↑	1.9	1.6	1.4	2.5	1.7
↗	1.4	0.58	1.8	1.6	1.6
↖	2.2	1.3	0.97	1.6	1.5
↙	2.5	2	1.3	0.9	1.5
↘	1.8	2.9	1.7	1.2	2.1
←	2.6	1.5	2.1	1.2	2.2

“Q-table”

Somme attendue des récompenses

On sélectionne l'action qui maximise la somme attendue des récompenses

$$a = \operatorname{argmax}_a Q(s, a)$$



**Observation**

**Actions**


	LEFT	DOWN	RIGHT	UP	STAY
→	1.7	1.4	2.6	0.86	1.5
↓	1.3	1.6	2.4	0.78	1.5
↑	1.9	1.6	1.4	2.5	1.7
↗	1.4	0.58	1.8	1.6	1.6
↖	2.2	1.3	0.97	1.6	1.5
↙	2.5	2	1.3	0.9	1.5
↘	1.8	2.9	1.7	1.2	2.1
←	2.6	1.5	2.1	1.2	2.2

“Q-table”

Somme attendue des récompenses

On sélectionne l'action qui maximise la somme attendue des récompenses

$$a = \operatorname{argmax}_a Q(s, a)$$



Observation

		Actions				
		LEFT	DOWN	RIGHT	UP	STAY
Observation	→	1.7	1.4	2.6	0.86	1.5
	↓	1.3	1.6	2.4	0.78	1.5
	↑	1.9	1.6	1.4	2.5	1.7
	↗	1.4	0.58	1.8	1.6	1.6
	↖	2.2	1.3	0.97	1.6	1.5
	↙	2.5	2	1.3	0.9	1.5
	↓	1.8	2.9	1.7	1.2	2.1
	←	2.6	1.5	2.1	1.2	2.2

“Q-table”

Somme attendue des récompenses

On sélectionne l'action qui maximise la somme attendue des récompenses

$$a = \operatorname{argmax}_a Q(s, a)$$




# Q-learning

“Equation de Bellman”


$$Q_{n+1}(s_t, a_t) = r_{t+1} + \gamma \max_a Q_n(s_{t+1}, a)$$

	LEFT	DOWN	RIGHT	UP	STAY
→	1.7	1.4	2.6	0.86	1.5
↘	1.3	1.6	2.4	0.78	1.5
↑	1.9	1.6	1.4	2.5	1.7
↗	1.4	0.58	1.8	1.6	1.6
↖	2.2	1.3	0.97	1.6	1.5
↙	2.5	2	1.3	0.9	1.5
↓	1.8	2.9	1.7	1.2	2.1
←	2.6	1.5	2.1	1.2	2.2





$$= \sum r_{t+1} + r_{t+2} + \dots + r_T$$




# Q-learning

		Actions				
		LEFT	DOWN	RIGHT	UP	STAY
States	→	1.7	1.4	2.6	0.86	1.5
	↘	1.3	1.6	2.4	0.78	1.5
	↑	1.9	1.6	1.4	2.5	1.7
	↗	1.4	0.58	1.8	1.6	1.6
	↖	2.2	1.3	0.97	1.6	1.5
	↙	2.5	2	1.3	0.9	1.5
	↓	1.8	2.9	1.7	1.2	2.1
	↔	2.6	1.5	2.1	1.2	2.2

$$a = \operatorname{argmax}_a Q(s_t, a)$$




# “Explorer” ou “exploiter” ?



# Explorer... parfois !

## Méthode “epsilon-greedy”



# Q-learning

States

	Actions				
	LEFT	DOWN	RIGHT	UP	STAY
→	1.7	1.4	2.6	0.86	1.5
↘	1.3	1.6	2.4	0.78	1.5
↑	1.9	1.6	1.4	2.5	1.7
↗	1.4	0.58	1.8	1.6	1.6
↖	2.2	1.3	0.97	1.6	1.5
↙	2.5	2	1.3	0.9	1.5
↓	1.8	2.9	1.7	1.2	2.1
←	2.6	1.5	2.1	1.2	2.2

$$a = \operatorname{argmax}_a Q(s_t, a)$$

+  
epsilon-greedy



Bellman Equation



# Q-learning



Fonctionne dans tous les cas !\*

States

	Actions				
	LEFT	DOWN	RIGHT	UP	STAY
→	1.7	1.4	2.6	0.86	1.5
↘	1.3	1.6	2.4	0.78	1.5
↑	1.9	1.6	1.4	2.5	1.7
↗	1.4	0.58	1.8	1.6	1.6
↖	2.2	1.3	0.97	1.6	1.5
↙	2.5	2	1.3	0.9	1.5
↓	1.8	2.9	1.7	1.2	2.1
←	2.6	1.5	2.1	1.2	2.2

$$a = \operatorname{argmax}_a Q(s_t, a)$$

+  
epsilon-greedy



Bellman Equation



# Q-learning



Fonctionne dans tous les cas !\*

States

	Actions				
	LEFT	DOWN	RIGHT	UP	STAY
→	1.7	1.4	2.6	0.86	1.5
↘	1.3	1.6	2.4	0.78	1.5
↑	1.9	1.6	1.4	2.5	1.7
↗	1.4	0.58	1.8	1.6	1.6
↖	2.2	1.3	0.97	1.6	1.5
↙	2.5	2	1.3	0.9	1.5
↓	1.8	2.9	1.7	1.2	2.1
←	2.6	1.5	2.1	1.2	2.2

$$a = \operatorname{argmax}_a Q(s_t, a)$$

+  
epsilon-greedy



Bellman Equation




\* peut prendre un temps infini 😅

## Bellman Equation

$$\underbrace{Q_{n+1}(s_t, a_t)}_{\text{new Q-value}} = r_{t+1} + \gamma \max_a Q_n(s_{t+1}, a)$$



	LEFT	DOWN	RIGHT	UP	STAY
→	1.7	1.4	2.6	0.86	1.5
↓	1.3	1.6	2.4	0.78	1.5
↑	1.9	1.6	1.4	2.5	1.7
↗	1.4	0.58	1.8	1.6	1.6
↖	2.2	1.3	0.97	1.6	1.5
↙	2.5	2	1.3	0.9	1.5
↓	1.8	2.9	1.7	1.2	2.1
↖	2.6	1.5	2.1	1.2	2.2



## Bellman Equation


$$Q_{n+1}(s_t, a_t) = r_{t+1} + \underbrace{\gamma \max_a Q_n(s_{t+1}, a)}_{\text{new Q-value}}$$

	LEFT	DOWN	RIGHT	UP	STAY
→	1.7	1.4	2.6	0.86	1.5
↓	1.3	1.6	2.4	0.78	1.5
↑	1.9	1.6	1.4	2.5	1.7
↗	1.4	0.58	1.8	1.6	1.6
↖	2.2	1.3	0.97	1.6	1.5
↙	2.5	2	1.3	0.9	1.5
↓	1.8	2.9	1.7	1.2	2.1
↖	2.6	1.5	2.1	1.2	2.2




## Equation de Bellman

$$Q_{n+1}(s_t, a_t) = r_{t+1} + \gamma \max_a Q_n(s_{t+1}, a)$$



**Discount factor**  
("facteur d'actualisation")

# Learning rate



$$\underbrace{Q_{n+1}(s_t, a_t)}_{\text{new Q-value}} = (1 - \alpha)Q_n(s_t, a_t) + \alpha(r_{t+1} + \gamma \max_a Q_n(s_{t+1}, a))$$

```
alpha = 0    # No update  
alpha = 1    # Full update  
alpha = 0.5 # Mean between old/new
```

Mise en pratique


# Introduction à Q-learning

[https://colab.research.google.com/drive/  
1KSei3ZdjNyyUYsPMqTyl4G2rm9v-P2E](https://colab.research.google.com/drive/1KSei3ZdjNyyUYsPMqTyl4G2rm9v-P2E)



## Q-table size

```
{'target_dir': 2, 'lidar': [0, 0, 0, 0, 1, 1, 1, 1]}
```




```
Dict(target_dir:Discrete(8), lidar:MultiBinary(8))
```

$$n_{\text{states}} = 8 * (2^{**8}) = 2048$$

## Q-table size

```
{'target_dir': 2, 'lidar': [0, 0, 0, 0, 1, 1, 1, 1]}
```



```
Dict(target_dir:Discrete(8), lidar:MultiBinary(8))
```

$$n_{\text{states}} = 8 * (2^{**8}) = 2048$$



# DQN


# Deep Q-Network

# Rappel : Q-learning

		Actions				
		LEFT	DOWN	RIGHT	UP	STAY
Observations	→	1.7	1.4	2.6	0.86	1.5
	↘	1.3	1.6	2.4	0.78	1.5
	↑	1.9	1.6	1.4	2.5	1.7
	↗	1.4	0.58	1.8	1.6	1.6
	↖	2.2	1.3	0.97	1.6	1.5
	↙	2.5	2	1.3	0.9	1.5
	↓	1.8	2.9	1.7	1.2	2.1
	←	2.6	1.5	2.1	1.2	2.2

On sélectionne l'action qui maximise la somme attendue des récompenses

$$a = \operatorname{argmax}_a Q(s, a)$$



# Rappel : Q-learning

“Equation de Bellman”

$$Q_{n+1}(s_t, a_t) = r_{t+1} + \gamma \max_a Q_n(s_{t+1}, a)$$

	LEFT	DOWN	RIGHT	UP	STAY
→	1.7	1.4	2.6	0.86	1.5
↘	1.3	1.6	2.4	0.78	1.5
↑	1.9	1.6	1.4	2.5	1.7
↗	1.4	0.58	1.8	1.6	1.6
↖	2.2	1.3	0.97	1.6	1.5
↙	2.5	2	1.3	0.9	1.5
↓	1.8	2.9	1.7	1.2	2.1
←	2.6	1.5	2.1	1.2	2.2



# Deep Q-network (DQN)

charge: 16.18%                ...

lidar\_left: 0            ...

lidar\_down: 1            ...

Q(Up) = 0.33    Q(Right) = 3.14    Q(Stay) = 2.71




## Deep Q-network (DQN)

charge: 16.18%    lidar\_left: 0    lidar\_down: 1

charge: 16.18%	lidar_left: 0	lidar_down: 1	
----------------	---------------	---------------	--

Q(Up) = 0.33    Q(Right) = 3.14    Q(Stay) = 2.71




## “Q-network”


$$\delta = Q_n(s_t, a_t) - (r_{t+1} + \gamma \max_a Q_n(s_{t+1}, a))$$

$\delta$  est l'erreur à minimiser


# Deep Q-network (DQN)



Couche dense (“fully connected”)



Couche à convolutions



Timestep: 1, action: left


Reward: 1.0

Timestep: 2, action: right

Reward: 1.0

Timestep: 3, action: left

Reward: 1.0



Timestep: 4, action: left

Reward: 1.0


Timestep: 5, action: left


Reward: 1.0

Timestep: 6, action: right

Reward: 0.0







# A vous de jouer !



Mise en pratique

## Deep Q-Network

<https://colab.research.google.com/drive/1R9X7vTaoTzDI1QGoArBMnMwgPB-BXK2C>



**Peut-on faire mieux ?**



# Les difficultés de l'AR

- Entrainement long
- Sensibilité des paramètres
- Manque de généralisation
- Dilemme exploration/exploitation

# Trucs et Astuces

- “Reward engineering”
  - Ajustement à la main des récompenses
  - Mais, risque d'introduire des biais !
- Entraîner dans différents environnements
  - Appeler plusieurs fois `trainer.train()`
- Faire varier le paramètre epsilon
  - Le faire descendre puis remonter de façon cyclique ?
- Entraîner plus n'améliore pas toujours les choses !
  - Évaluer régulièrement, puis garder le meilleur agent ?

# Les difficultés de l'AR

- Entrainement long
- Sensibilité des paramètres
- Manque de généralisation
- Dilemme exploration/exploitation

Apprentissage Priorisé !



# Prioritized Experience Replay


(“Rejeu d'Expériences Priorisées”)


Schaul et al. 2015


<https://arxiv.org/abs/1511.05952>


# But : **apprendre plus efficacement**



















“Toutes les expériences ne sont pas égales”



A Katharopoulos






99% du replay buffer







Les 1% !

Sample in smarter way!









<https://www.theguardian.com/news/datablog/2013/may/31/times-tables-hardest-easiest-children>







Erreurs faibles







Erreurs élevées !





Trop de biais!





# Réduisons le biais...

Avant

$\alpha$

Débiaisage





# Réduisons le biais...

$\alpha$

Débiaisage

Après





# Réduisons le biais...

$\alpha$

Débiaisage

$$P(i) = \frac{p_i^\alpha}{\sum_k p_k^\alpha}$$

Après





# Réduisons le biais...

$\alpha$

Débiaisage

$\beta$

Correction de l'échantillonnage




# Réduisons le biais...

Plus d'effet

$\alpha$


Débiaisage



$\beta$

Correction de l'échantillonnage

Moins d'effet




# Réduisons le biais...

Plus d'effet

$\alpha$


Débiaisage  
= 0.6







$\beta$


Correction de l'échantillonnage  
= 0.4

Moins d'effet






— human — uniform — rank-based — proportional



# Rainbow DQN



# A vous de jouer !



Mise en pratique

# Apprentissage Priorisé

[https://colab.research.google.com/drive/  
156LKB0aPnKab1loUjeEKwV3VTmL7fguW](https://colab.research.google.com/drive/156LKB0aPnKab1loUjeEKwV3VTmL7fguW)

